

Mohamad Liyaghi

Software Engineer · AI

Tehran, Iran · +98 991 951 4231 · liaghimohamad69@gmail.com · mohamad-liyaghi.ir · github.com/mohamad-liyaghi · linkedin.com/in/mohamad-liyaghi · medium.com/@el_mohamad

SUMMARY

Software engineer with 5+ years building scalable, high-throughput backend systems and shipping production AI features. Full-stack across Python, TypeScript and Go with a deep backend focus — distributed services, async pipelines and real-time systems — plus hands-on LLM work: agents, RAG, guardrails and cost engineering. Comfortable end to end, from API design and data modelling to containers, cloud and observability. I write regularly about the realities of running AI in production.

SKILLS

Languages	Python, TypeScript, JavaScript, Go
Backend	Django, Django REST Framework, FastAPI, Celery, Express, Go Fiber
Frontend	React, Next.js, TypeScript, Tailwind CSS, Remotion (programmatic video)
AI / LLM	LLM integration, RAG, AI agents, LangChain, guardrails, cost engineering, vector search
Data & Search	PostgreSQL, PostGIS, MongoDB, Redis, MySQL, Qdrant, Milvus, Elasticsearch
Async / Realtime	RabbitMQ, Celery, Socket.IO, WebSockets, asyncio, microservices
DevOps / Cloud	Docker, Kubernetes, GitHub Actions, Nginx, AWS, Unicorn, Uvicorn
Observability	Sentry, OpenTelemetry, Prometheus, Grafana, Loki, Elastic APM
Testing	Pytest, unittest

EXPERIENCE

Software Engineer · Team Lead — Zebracat

Sep 2024 – Present · Remote (Berlin, DE)

AI text-to-video platform (zebracat.ai)

- Lead backend development for an AI text-to-video platform — microservices that fan out heavy generative jobs and stream real-time progress back to users.
- Design async pipelines (queues, workers, back-pressure) that keep latency and cost predictable as request volume scales.
- Integrate and productionize LLM features — agents, retrieval and guardrails — with the observability and cost controls that keep them reliable.
- Lead the team day to day: scoping, code review and on-time delivery while keeping services stable as the product moves fast.

Full-Stack Developer — Freelance

Jun 2021 – Sep 2024 · Remote

- Delivered end-to-end web products (SaaS and consumer) from requirements to production, owning both frontend and backend.
- Built scalable Django / FastAPI backends and responsive React frontends; handled deployment, monitoring and third-party integrations.
- Shipped reliable, secure systems under tight client deadlines.

SELECTED PROJECTS

FoodAnywhere — Open-source food-delivery backend with geospatial search and worker-offloaded tasks.

Django, PostGIS, Redis, Celery, Prometheus, Grafana, Loki, Docker, Kubernetes.

FastQuora — Q&A platform with full-text search, caching and end-to-end tests. FastAPI, PostgreSQL, Elasticsearch, Redis, Docker, Pytest.

fast-commerce — Async e-commerce backend focused on throughput and clean API design. FastAPI, PostgreSQL, Redis, Pytest, Kubernetes.

AcademyMaster — Academy-management API with layered permissions and background jobs. Django, DRF, Celery, PostgreSQL, Docker.

SELECTED WRITING — [MEDIUM.COM/@EL_MOHAMAD](https://medium.com/@el_mohamad)

- Stop Treating LLM Calls as Free: How Real Cost Engineering for Agents Actually Works (2026)
- Stop Trusting the Embedding: How Real RAG Pipelines Actually Work (2026)
- Stop Trusting the Model: How Real LLM Guardrails Actually Work (2026)
- Stop Putting Everything in One Prompt: How Real Research Agents Actually Work (2026)

LANGUAGES

Persian — Native

English — Professional (C1)